

Large-Scale Distributed Bayesian Matrix Factorization using Stochastic Gradient MCMC

Sungjin Ahn, Anoop Korattikara, Nathan Liu,
Suju Rajan, Max Welling


May 30, 2017

Bayesian MF Literature



Recap: Bayesian PMF

PMF:




- ▶ Rating Matrix: $R \in \mathbb{R}^{L \times M}$ ($\#$ users = L , $\#$ items = M)
- ▶ User latent features: $U \in \mathbb{R}^{D \times L}$
- ▶ Movie latent features: $V \in \mathbb{R}^{D \times M}$

Recap: Bayesian PMF

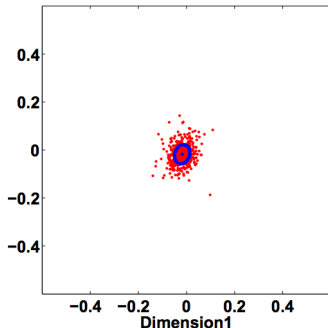
PMF \rightarrow Bayesian PMF

- ▶ In reality, R is very sparse!
- ▶ MLE estimation for U and V lead to severe over-fitting
- ▶ We want to model the uncertainty in the items

Movie X (5 ratings)




Movie Y (142 ratings)



Recap: Bayesian PMF

Model

- ▶ Introduce priors for the parameters
- ▶ Allow model complexity to be controlled automatically



$$p(\mathbf{U}|\mu_U, \Lambda_U) = \prod_{i=1}^L \mathcal{N}(U_i|\mu_U, \Lambda_U^{-1}),$$
$$p(\mathbf{V}|\mu_V, \Lambda_V) = \prod_{j=1}^M \mathcal{N}(V_j|\mu_V, \Lambda_V^{-1}).$$

$$p(\Theta_U|\Theta_0) = \mathcal{N}(\mu_U|\mu_0, (\beta_0\Lambda_U)^{-1})\mathcal{W}(\Lambda_U|W_0, \nu_0),$$

$$p(\Theta_V|\Theta_0) = \mathcal{N}(\mu_V|\mu_0, (\beta_0\Lambda_V)^{-1})\mathcal{W}(\Lambda_V|W_0, \nu_0),$$

Recap: Bayesian PMF

Inference

- predictive distribution:

$$p(R_{ij}^* | \mathbf{R}, \Theta_0) = \int \int p(R_{ij}^* | \underline{U_i, V_j}) p(\mathbf{U}, \mathbf{V} | \mathbf{R}, \Theta_U, \Theta_V) p(\Theta_U, \Theta_V | \Theta_0) d\{\mathbf{U}, \mathbf{V}\} d\{\Theta_U, \Theta_V\}$$

- MCMC

$$p(R_{ij}^* | \mathbf{R}, \Theta_0) \approx \frac{1}{T} \sum_{t=1}^T p(R_{ij}^* | U_i^{(t)}, V_j^{(t)}).$$

where $\{U_i^t, V_j^t\}$ is generated by a markov-chain whose stationary distribution is the posterior distribution over the model parameters:

$$\begin{aligned} & p(\mathbf{U}, \mathbf{V}, \Theta_U, \Theta_V | \mathbf{R}, \Theta_0) \\ &= p(\mathbf{U}, \mathbf{V} | \mathbf{R}, \Theta_U, \Theta_V) p(\Theta_U, \Theta_V | \Theta_0) \end{aligned}$$

Recap: Bayesian PMF

Inference

Algorithm 1 Gibbs Sampling for BPMF

```
1: Initialize model parameters  $\mathbf{U}^{(1)}, \mathbf{V}^{(1)}$ 
2: for  $t = 1 : T$  do
3:   // Sample hyperparameters
    $\Theta_U^{(t)} \sim p(\Theta_U | \mathbf{U}^{(t)}, \Theta_0), \quad \Theta_V^{(t)} \sim p(\Theta_V | \mathbf{V}^{(t)}, \Theta_0)$ 
4:   for  $i = 1 : L$  in parallel do
5:      $U_i^{(t+1)} \sim p(U_i | \mathbf{R}, \mathbf{V}^{(t)}, \Theta_U^{(t)})$  // sample user features
6:   end for
7:   for  $j = 1 : M$  in parallel do
8:      $V_j^{(t+1)} \sim p(V_j | \mathbf{R}, \mathbf{U}^{(t)}, \Theta_V^{(t)})$  // sample item features
9:   end for
10: end for
```

Recap: Bayesian PMF

Inference

- ▶ Sampling from conditionals for U_i, V_j in Gibbs, involve $O(D^3)$ computation per iteration (inverting $D \times D$ precision matrix)

$$\begin{aligned} p(U_i | R, V, \Theta_U, \alpha) &= \mathcal{N}(U_i | \mu_i^*, [\Lambda_i^*]^{-1}) \\ &\sim \prod_{j=1}^M \left[\mathcal{N}(R_{ij} | U_i^T V_j, \alpha^{-1}) \right]^{I_{ij}} p(U_i | \mu_U, \Lambda_U), \end{aligned}$$

where

$$\begin{aligned} \Lambda_i^* &= \Lambda_U + \alpha \sum_{j=1}^M [V_j V_j^T]^{I_{ij}} \\ \mu_i^* &= [\Lambda_i^*]^{-1} \left(\alpha \sum_{j=1}^M [V_j R_{ij}]^{I_{ij}} + \Lambda_U \mu_U \right). \end{aligned}$$

- ▶ Each iteration of MCMC requires computations over the entire dataset

Langevin Dynamics

$$\theta_{t+1} \leftarrow \theta_t + \frac{\epsilon_t}{2} \left(\nabla_{\theta_t} \log p(\theta_t) + \sum_{x \in X} \nabla_{\theta_t} \log p(x|\theta) \right) + \mathcal{N}(0, \epsilon_t)$$

- ▶ Converges to true posterior distribution if step size ϵ_t is annealed to 0 at a rate satisfying:
 $\sum_{t=1}^{\infty} \epsilon_t = \infty, \quad \sum_{t=1}^{\infty} \epsilon_t^2 < \infty$
- ▶ No accept-reject tests since acceptance rate $\rightarrow 1$ as $\epsilon_t \rightarrow 0$.
- ▶ Requires $O(N)$ computation per iteration.

Stochastic Gradient Langevin Dynamics (SGLD)

[Welling and Teh, 2011]

$$\theta_{t+1} \leftarrow \theta_t + \frac{\epsilon_t}{2} (\text{Stochastic Gradient}) + \mathcal{N}(0, \epsilon_t)$$

$$\theta_{t+1} \leftarrow \theta_t + \frac{\epsilon_t}{2} \left(\nabla_{\theta_t} \log p(\theta_t) + \frac{N}{m} \sum_{x \in \mathcal{M}} \nabla_{\theta_t} \log p(x|\theta) \right) + \mathcal{N}(0, \epsilon_t)$$

- $O(m \ll N)$ computation per iteration.

Bayesian PMF using SGLD

Model

$$p(\mathbf{U}|\Lambda_U) = \prod_{i=1}^L \mathcal{N}(U_i|0, \Lambda_U^{-1}),$$

$$p(\mathbf{V}|\Lambda_V) = \prod_{j=1}^M \mathcal{N}(V_j|0, \Lambda_V^{-1}).$$

$$\lambda_{U_d}, \lambda_{V_d} \sim \text{Gamma}(\alpha_0, \beta_0).$$

- ▶ Slightly simplified model
- ▶ Λ_U and Λ_V are D -dimensional diagonal matrices with diagonal elements λ_{U_d} and λ_{V_d} respectively
- ▶ Each latent vector can be updated in $O(D)$ instead of $O(D^2)$ in the full-covariance case

Bayesian PMF using SGLD

Inference

- Recall the full posterior:

$$\begin{aligned} p(\mathbf{U}, \mathbf{V}, \Theta_U, \Theta_V | \mathbf{R}, \Theta_0) \\ = p(\mathbf{U}, \mathbf{V} | \mathbf{R}, \Theta_U, \Theta_V) p(\Theta_U, \Theta_V | \Theta_0) \end{aligned}$$

- Alternate b/w steps:
 - Step 1: Sample from $p(U, V | R, \Theta_U, \Theta_V)$ using SGLD
 - Step 2: Sample from $p(\Theta_U, \Theta_V | \Theta_0)$ using Gibbs

Bayesian PMF using SGLD

Step 1: Sample from $p(U, V|R, \Theta_U, \Theta_V)$ using SGLD

- ▶ Suppose rating matrix is: $\mathcal{X} = \{x_n = (p_n, q_n, r_n)\}_{n=1}^N$
- ▶ Gradient of log-posterior w.r.t U_i

$$G(\mathcal{X}) = \sum_{n=1}^N g_n(U_i; \mathcal{X}) - \Lambda_U U_i$$

$$g_n(U_i; \mathcal{X}) = \mathbb{I}[p_n = i | \mathcal{X}] (r_n - U_{p_n}^\top V_{q_n}) V_{q_n}$$

- ▶ To use SGLD, need unbiased estimate of this gradient from a mini-batch $\mathcal{M} = \{(p_n, q_n, r_n)\}_{n=1}^m$ of m tuples of \mathcal{X}

$$G_1(\mathcal{M}) = N \bar{g}(U_i; \mathcal{M}) - \Lambda_U U_i$$

$$\bar{g}(U_i; \mathcal{M}) = \frac{1}{m} \sum_{n=1}^m g_n(U_i; \mathcal{M})$$

Bayesian PMF using SGLD

Step 1: Sample from $p(U, V|R, \Theta_U, \Theta_V)$ using SGLD

- Unbiased estimate of the gradient which needs only updating users in the mini-batch \mathcal{M}

$$G_3(\mathcal{M}) = N\bar{g}(U_i; \mathcal{M}) - \mathbb{I}[i \in \mathcal{M}_p] h_{i*}^{-1} \Lambda_U U_i.$$

$$h_{i*} = 1 - \left(1 - \frac{N_{i*}}{N}\right)^m$$

$$N_{i*} = \sum_{n=1}^N \mathbb{I}[p_n = i | \mathcal{X}]$$

- Overall update for U_i and V_j (s is a worker)

$$U_{i,t+1} \leftarrow U_{i,t} + \frac{\epsilon_t}{2} \left\{ \frac{N^{(s)}}{v^{(s)}} \bar{g}(U_{i,t}; \mathcal{M}_t^{(s)}) - \frac{\Lambda_U U_{i,t}}{\bar{h}_{i*}} \right\} + \nu_t \quad (29)$$

$$V_{j,t+1} \leftarrow V_{j,t} + \frac{\epsilon_t}{2} \left\{ \frac{N^{(s)}}{v^{(s)}} \bar{g}(V_{j,t}; \mathcal{M}_t^{(s)}) - \frac{\Lambda_V V_{j,t}}{\bar{h}_{*j}} \right\} + \nu_t. \quad (30)$$

Bayesian PMF using SGLD

Mini-batch for Step 1

| | | |
|---|---|---|
| 2 | 1 | 3 |
| 2 | 2 | 1 |
| 3 | 2 | 2 |
| 1 | 3 | 5 |
| 2 | 4 | 2 |
| 3 | 3 | 4 |
| 4 | 1 | 2 |
| 5 | 1 | 4 |
| 6 | 2 | 1 |
| 4 | 4 | 1 |
| 5 | 3 | 5 |
| 6 | 4 | 3 |

Bayesian PMF using SGLD

Step 2: Sample from $p(\Theta_U, \Theta_V | \Theta_0)$ using Gibbs

- This is easy because of conjugacy:


$$\lambda_{U_d} | \mathbf{U}, \mathbf{V} \sim \text{Gamma} \left(\alpha_0 + \frac{L}{2}, \beta_0 + \frac{1}{2} \sum_{i=1}^L U_{di}^2 \right), \quad (24)$$

$$\lambda_{V_d} | \mathbf{U}, \mathbf{V} \sim \text{Gamma} \left(\alpha_0 + \frac{M}{2}, \beta_0 + \frac{1}{2} \sum_{i=1}^M V_{di}^2 \right). \quad (25)$$

Distributed SGLD


Parameter Server Architecture

- ▶ Worker s perform Step 1 - Sample from $p(U_s, V_s | R, \Theta_{U_s}, \Theta_{V_s})$
- ▶ Master performs Step 2 - Sample from $p(\Theta_U, \Theta_V | \Theta_0)$



Distributed SGLD

Distribution of data (rating matrix R) and U, V among workers



- ▶ Rating matrix R is partitioned into mutually-exclusive blocks B_{ij}
- ▶ Run 2 independent chains with parameters $\{U^a, V^a\}$ and $\{U^b, V^b\}$

DSGLD: Algorithm (parameter server)

Algorithm 2 DSGLD at parameter server

```
1: Initialize model parameters of each chain  $\{\mathbf{U}_1^c, \mathbf{V}_1^c, \Lambda_1^c\}_{c=1}^C$ ,  
   step sizes  $\{\epsilon_t\}$   
2: for each chain  $c$  parallel do  
3:   for  $t=1:\text{max\_iter}$  do  
4:      $\mathcal{B}_c \leftarrow \text{GET\_ORTHO\_BLOCK\_GROUP}(c, t)$   
5:     for worker  $s \in \mathcal{B}_c$  do  
6:        $\mathbf{U}_{t+1}^{(c,s)}, \mathbf{V}_{t+1}^{(c,s)} \leftarrow \text{WKR\_ROUND}(\mathbf{U}_t^{(c,s)}, \mathbf{V}_t^{(c,s)}, \Lambda_t^{(c)}, \epsilon_t)$   
7:     end for  
8:     if not burn-in then  
9:       Store  $\mathbf{U}_{t+1}^{(c)}, \mathbf{V}_{t+1}^{(c)}$  as a sample of chain  $c$   
10:      Sample  $\Lambda_{t+1}^{(c)} | \mathbf{U}_{t+1}^{(c)}, \mathbf{V}_{t+1}^{(c)}$  using Eqn. (24) and (25)  
11:    end if  
12:  end for  
13: end for
```

DSGLD: Algorithm (worker)

Algorithm 3 DSGLD at worker s

```
1: Initialize  $\bar{h}_{i*}, \bar{h}_{*j}$ , round length  $\gamma$ , mini-batch size  $m$ 
2: function WKR_ROUND( $\mathbf{U}^{(c,s)}, \mathbf{V}^{(c,s)}, \Lambda^{(c)}, \epsilon_t$ )
3:   for  $t = 1 : \gamma$  do
4:     Sample a mini-batch  $\mathcal{M}_t$  from  $\mathcal{X}^{(s)}$ 
5:     for each user  $i$  and item  $j$  in  $\mathcal{M}_t$  parallel do
6:       Update  $U_i, V_j$  using Eqn. (29) and (30)
7:     end for
8:   end for
9:   Send updated  $\mathbf{U}^{(c,s)}$  and  $\mathbf{V}^{(c,s)}$  to the parameter server
10: end function
```

DSGLD: Block Split Schemes

| | |
|---|---|
| 1 | 2 |
| 3 | 4 |

(a) square

| |
|---|
| 1 |
| 2 |
| 3 |
| 4 |

(b) column

Two levels of parallelism

- ▶ (a) within-chain (both X and $\{U, V\}$ are partitioned)
- ▶ (b) between-chain (X and U are partitioned, entire V is transmitted to parameter server)

Experiments

| | Optimization | MCMC |
|----------------|--------------|-------------|
| Single Machine | SGD | SGLD, Gibbs |
| Distributed | DSGD | DSGLD |

| Dataset | # users | # items | # ratings |
|---------|---------|---------|-----------|
| Netflix | 480K | 18K | 100M |
| Yahoo | 1.8M | 136K | 700M |

Experiments





Figure: Netflix dataset ($D = 30$)

- Single node, multi core

Experiments







Figure: Yahoo Music Rating dataset ($D = 30$)

- 16 nodes

Experiments




(a) DSGLD-C on the Netflix dataset




(b) DSGLD-S on the Yahoo dataset

Figure: The effect of the number of chains, number of workers, and block split


Experiments



(a) RMSE on Netflix



(b) RMSE on Yahoo



(c) Required time per sample

Figure: The effect of the latent feature dimension

Conclusion

- ▶ Paper combines ideas from DSGLD to scale Bayesian MF
- ▶ Avoids $O(N)$ computation over the entire dataset
- ▶ Using multiple chains and multiple workers in parallel gives the method better and faster mixing
 - ▶ More chains help explore a broader space
 - ▶ Updating orthogonal blocks in parallel helps chain mix faster
 - ▶ Averaging predictions from independent chains gives better generalization
- ▶ Limitations: Slow worker, Not fully model parallel (model parallelism at worker)