


Cover Trees for Nearest Neighbor

Alina Beygelzimer, Sham Kakade, John Langford

(Presented by: Parameswaran Raman, params@ucsc.edu)

April 11, 2016

Nearest Neighbor Problem



Given a set S of n -points in some metric space (X, d) , pre-process S s.t given a query point $q \in X$, one can efficiently find a point $p \in S$ which *minimizes* $d(p, q)$.

Comparison of Methods and Challenges

Brute Force



- ▶ requires no pre-processing
- ▶ query time is $O(n)$, space is $O(n)$

Comparison of Methods and Challenges

K-d Tree [FBL77]

Data points in 2D: (x_1, x_2)



$(1,9), (2,3), (4,1), (3,7), (5,4), (6,8), (7,2), (8,8), (7,9), (9,6)$



Comparison of Methods and Challenges

K-d Tree [FBL77]


Query point: (7, 4)



- ▶ works only for low dimensions (≈ 10)
- ▶ inexact NN

Comparison of Methods and Challenges

Ball Tree / Metric Tree [Omo87][Uhl91]



- ▶ works on moderately high dimensions (better than K-d trees)
- ▶ more generally applicable than K-d trees

Comparison of Methods and Challenges

Methods that assume structure (intrinsic dimensionality)

[KR02][KL04b]

- ▶ Karger and Ruth - [KR02] defines *expansion constant* - a measure of intrinsic dimensionality
- ▶ Navigating Nets by [KL04b] uses a similar measure
- ▶ Drawback: Although query times are good, space requirements are exponential in d (# of dimensions)

What are the highlights of the paper?


	Brute Force	K-d Tree	Ball Tree	Cover Tree
Const Time	$O(n)$	$O(n \log n)$	$O(n^2)$	$O(c^6 n \log n)$
Const Space	$O(n)$	$O(n)$	$O(n)$	$O(n)$
Insert/Remove	$O(n)$	$O(\log n)$	$O(n)$	$O(c^6 \log n)$
Query Time	$O(n)$	$O(\log n)$	$O(n)$	$O(c^{12} \log n)$

- ▶ Linear space requirement independent of metric structure and dimensionality
- ▶ Query time as good as or better than other methods
- ▶ Theoretical guarantees
- ▶ Efficient implementations available

Assumption about the structure: Expansion Constant

$B(p, r) = \{q \in S : d(q, p) < r\}$ = points within distance r of p .

$$\exists c: \forall p \in S, r > 0: c|B(p, r)| \geq |B(p, 2r)|$$




c = Expansion Constant of the set of points S

Assumption about the structure: Expansion Constant

If S (set of points in the metric space X) is arranged uniformly on some surface of dimension d , then:

$$c \approx 2^d$$

- ▶ Eg: grid of equally spaced points in 2-D will have $c = 4$




- ▶ Eg: grid of equally spaced points in 3-D will have $c = 8$
- ▶ Other interesting cases

What is a Cover Tree?


- ▶ Leveled tree satisfying invariants
- ▶ One point per node
- ▶ Lemma: for all nodes, number of children $\leq c^4$
- ▶ Lemma: maximum depth $= O(c^2 \log n)$

What is a Cover Tree?




- (Nesting invariant): $C_i \subseteq C_{i-1}$ (C_i = nodes at level i)

What is a Cover Tree?




- (Covering Tree invariant): $\forall p \in C_{i-1}, \exists q \in C_i: d(p,q) \leq 2^i$ (p is a child of q)

What is a Cover Tree?




- (Separation invariant): If $p, q \in C_i$, then $d(p, q) \geq 2^i$

What is a Cover Tree?




- Explicit Representation = Implicit Representation with duplicates removed

How to use the Cover Tree for NN Query?




- ▶ Start with upper bound at root
- ▶ Descend tree, maintain a "cover set"

How to use the Cover Tree for NN Query?



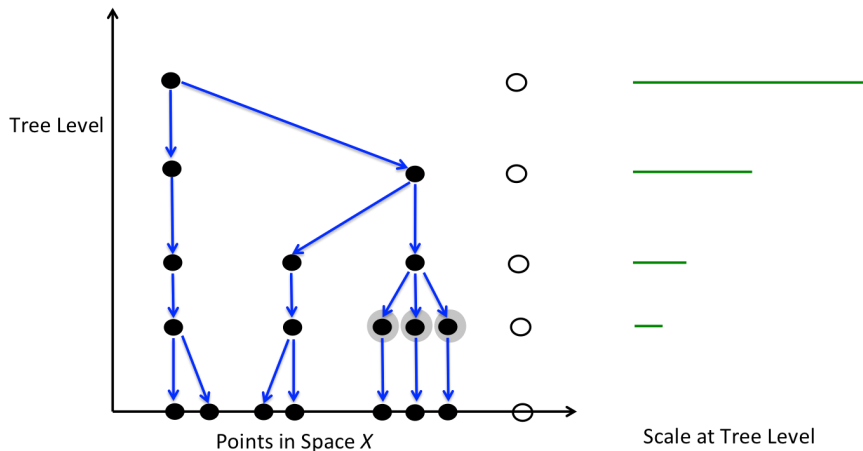
- ▶ Start with upper bound at root
- ▶ Descend tree, maintain a "cover set"

How to use the Cover Tree for NN Query?




- ▶ Start with upper bound at root
- ▶ Descend tree, maintain a "cover set"

How to use the Cover Tree for NN Query?



- ▶ Start with upper bound at root
- ▶ Descend tree, maintain a "cover set"

How to use the Cover Tree for NN Query?




- ▶ Start with upper bound at root
- ▶ Descend tree, maintain a "cover set"


How to make a Cover Tree?

- ▶ Single Point Insertion - $O(c^6 \log n)$
- ▶ Batch/Lazy Construction

Empirical Study





Speedup over Brute Force



Speedup over sb(S) data structure [Ken Clarkson]

- Note: mnist is the largest dataset ($n=60,000$, $d=784$)

Effects of expansion constant



Summary

Cover Trees are:

- ▶ space-wise efficient than all other discussed methods
- ▶ time-wise on par with other discussed methods
- ▶ theoretically more elegant than K-d Trees / Metric Trees
- ▶ empirically not fully convincing yet - experiments need to be done on large-scale datasets comparing against Ball Trees and Navigation Nets
- ▶ Code: http://hunch.net/~jl/projects/cover_tree/cover_tree.html

References

- [1] Alina Beygelzimer, Sham Kakade, and John Langford. Cover trees for nearest neighbor. In *Proceedings of the 23rd international conference on Machine learning*, pages 97–104. ACM, 2006.
- [2] Kenneth L Clarkson. Nearest neighbor queries in metric spaces. *Discrete & Computational Geometry*, 22(1):63–93, 1999.
- [3] Jerome H Friedman, Jon Louis Bentley, and Raphael Ari Finkel. An algorithm for finding best matches in logarithmic expected time. *ACM Transactions on Mathematical Software (TOMS)*, 3(3):209–226, 1977.
- [4] David R Karger and Matthias Ruhl. Finding nearest neighbors in growth-restricted metrics. In *Proceedings of the thirty-fourth annual ACM symposium on Theory of computing*, pages 741–750. ACM, 2002.
- [5] Robert Krauthgamer and James R Lee. Navigating nets: simple algorithms for proximity search. In *Proceedings of the fifteenth annual ACM-SIAM symposium on Discrete algorithms*, pages 798–807. Society for Industrial and Applied Mathematics, 2004.
- [6] Stephen M Omohundro. *Efficient algorithms with neural network behavior*. Department of Computer Science, University of Illinois at Urbana-Champaign, 1987.
- [7] Jeffrey K Uhlmann. Satisfying general proximity/similarity queries with metric trees. *Information processing letters*, 40(4):175–179, 1991.